

The FOMC Window: What Is Identified and What Is Not

Progress memo

11 August 2026

Summary

I set out to test whether the Fed moves long-term rates by revealing information about the equilibrium real rate (r^*), using the FOMC’s published longer-run projection. That design does not work, for reasons I can now state precisely, and establishing why has produced three results worth reporting.

The headline finding is a timing fact that reframes the literature: **roughly 81% of the in-window decline in long yields occurs before 2012, and the dot plot begins in 2012.** The period in which the Fed’s endpoint view is observable is the period in which the phenomenon being explained is essentially absent. Every projection-based test of the information channel — including the one in the published paper — is therefore run where there is little to explain.

Alongside that I have a finished empirical result reconciling two contradictory working papers, and a methodological result showing why the standard decomposition cannot settle the question either way.

1 Where the project stands

The motivating fact is Hillenbrand (2025, *RFS*): the secular decline in long-term Treasury yields accumulates almost entirely in three-day windows around FOMC announcements. I replicate it — June 1989 to June 2021, the 10-year yield falls 6.99 pp, of which 6.16 pp (88%) accrues in windows and 0.82 pp outside — and extend the sample through July 2026 (Figure 1).

The result is best stated as an absence. The per-meeting effect is only -2.5 bp. What is remarkable is that non-window days drift -0.012 bp/day ($t = -0.18$) across 7,035 days. A randomisation test placing 324 randomly timed three-day blocks on non-FOMC days puts the observed -6.16 pp at the 0th percentile of 5,000 draws ($p < 0.0002$), so the concentration is not chance.

The natural question — and the one the literature is fighting over — is which component of the yield moves: expected future short rates, or the term premium.

2 Result 1: the literature’s disagreement is a sample-definition artifact

Two circulating papers give opposite answers. Hofmann, Li & Wu (BIS WP 1252) decompose the whole window and conclude it is risk-neutral. Pan & Peng (SSRN 4764451) decompose the pre-announcement day and conclude it is term premium.

I reproduce Pan & Peng’s estimates on their specification — scheduled meetings only, September 1994 onward — to three coefficients:




Figure 1: Replication and extension. Panels A–B reproduce the paper’s sample (June 1989 – June 2021): the 10-year yield falls 6.99 pp, 6.16 pp of it inside three-day FOMC windows and 0.82 pp outside. Panels C–D extend to 31 July 2026. The in-window series continues down to –7.24 pp while the outside series turns sharply positive to +3.56 pp — the entire 2022–2026 tightening in long rates happened on non-FOMC days. GSW zero-coupon yields; all 373 events including conference calls and unscheduled actions (11.8% of trading days).

Day $t - 1$, 10-year	Mine	Theirs
Yield	–0.785 ($t - 2.39$)	–0.79 ($t - 2.39$)
Term premium	–0.716 ($t - 2.36$)	–0.71 ($t - 2.34$)
Expected short rate	–0.070 ($t - 0.31$)	–0.08 ($t - 0.26$)

and then show the attribution is not robust. The term-premium share of the day- $(t-1)$ decline moves from **91%** (scheduled meetings, 1994 onward) to **46%** (all events, 1989 onward). The mechanism is identified: the 73 conference calls and unscheduled actions that a scheduled-only filter removes contribute –1.78 pp of risk-neutral decline on day $t - 1$ alone, about –2.5 bp per event against a scheduled-meeting average of –0.07 bp. Emergency intermeeting actions are preceded by days of expectations repricing; scheduled announcements are not.

Neither paper discusses the event list. This is a finished result requiring no model to be correct, and it is the most immediately publishable thing I have.

3 Result 2: the standard decomposition cannot settle the question

In a spanned affine model the state is recoverable from yields, so the expectations component is a fixed linear function of the curve: $\Delta r n_t^{(n)} = c'_n \Delta y_t^N$ with $c'_n = b'_n \mathbf{B}^{-1}$. Regressing the daily change in the ACM 10-year risk-neutral rate on six yield changes gives $R^2 = 1.00000$.

The consequence for event studies is that component responses are determined by the maturity profile of the yield response: $\beta^{rn} = c' \beta^N$. Applied to my estimated profile this predicts +2.5 and +29.3 bp against directly estimated +2.6 and +29.2. The decomposition adds no event-specific information; it applies historically estimated persistence weights to a profile already in hand.

Those weights are also known to be wrong for this application. Since the VAR is stationary, $b_n \rightarrow 0$ and the expectations component is asymptotically constant, so genuine endpoint variation is mechanically booked as term premium. Empirically $c' \mathbf{1} = 0.529$: a permanent, credible parallel shift — 100% expectations by construction — is reported as 53% expectations and 47% term premium.

Attribution. The algebra is a corollary of Joslin, Singleton & Zhu (2011) and appears as a criticism in Joslin, Priebisch & Singleton (2014). The stationarity bias is Bauer & Rudebusch (2020, *AER*), who report 75% versus 35–37% for the secular decline, and Bauer & Rudebusch (2014), who report the $\approx 50/50$ event-window figure. My contribution is the explicit filter for ACM and the event-study corollary, not the theorem.

4 Result 3: the effect and the measurement do not overlap in time

Splitting the in-window drift by period:

Period	Meetings	bp/meeting	t	Cumulative
Jun 1989 – Dec 2011 (pre-SEP)	245	−2.46	−3.06	− 5.85 pp
Jan 2012 – Mar 2022 (dot falling 1.78 pp)	85	−0.40	−0.41	−0.23 pp
Mar 2022 – Jul 2026 (dot rising 0.74 pp)	35	−3.49	−1.53	−1.16 pp

Cumulative column sums daily in-window changes; days shared by adjacent events are counted once.

Two things follow (Figure 2, Panel C). First, the phenomenon is overwhelmingly pre-2012, so projection-based tests are estimated where the effect is weakest. Second, during the decade when the FOMC cut its own longer-run rate by 1.78 pp — the episode in which an information channel should be most visible — the in-window contribution is indistinguishable from zero.

The 2022–2026 extension sharpens this. The Fed raised the target 525 bp and revised its longer-run dot *up* 0.74 pp. The 10-year rose 2.60 pp. In FOMC windows it fell 1.20 pp; the entire increase occurred outside. No paper in this literature covers this period — Hillenbrand ends June 2021, Hofmann–Li–Wu 2023, Adrian et al. November 2021.

5 What I have set aside, and why

- **The Fed–market wedge** (Figure 2). The market leg requires an endpoint measure. The ACM 9y1y risk-neutral forward has $\beta = 0.287$ on the one-year yield with $R^2 = 0.83$ and a




Figure 2: Why the r^* design fails. *Panel A:* the intended treatment — the FOMC's own longer-run real rate (SEP longer-run projection minus the 2% target). It exists only from 2012. *Panel B:* the intended market counterpart. The ACM-implied 9y1y risk-neutral forward loads on the current one-year yield with $\beta = 0.287$, $R^2 = 0.83$ and a level correlation of $+0.984$; it is a damped copy of the policy cycle, not an endpoint estimate (HLW, by contrast, is nearly orthogonal to the cycle). *Panel C:* the timing problem — 5.85 pp of the 7.24 pp cumulative in-window decline, roughly 81%, is complete before the dot plot begins.

level correlation of 0.984; it is a damped copy of the policy cycle, not an r^* estimate. Survey alternatives correlate -0.12 with it and barely move.

- **A credibility asymmetry between dovish and hawkish signals.** With 18 upward revisions I have 21–31% power. More importantly, Hillenbrand's Table 5 and Figure 11 already split by direction, and the ratchet mechanism is Rungcharoenkitkul, Borio & Disyatat (BIS WP 817).
- **Expectations versus term premium as a standalone question.** Contested by the two papers above, and unidentifiable for the reason in §3.

6 Path forward

Immediate (four weeks)

(a) **The issuance test.** Lou, Pintér, Üslü & Walker (BIS WP 1232) document a supply-driven pre-announcement drift in UK gilts and claim Hillenbrand's effect concentrates in FOMC windows *without* long-dated Treasury issuance. If correct, the mechanism is dealer inventory rather than monetary policy. Auction dates are public; this is a few days' work and it is decisive either way.

(b) **Draft the reconciliation note** (§2 plus the timing fact in §4). Self-contained, model-free, and it resolves a live disagreement.

Medium term (one term)

(c) **Permanent versus transitory by day type.** Window and non-window days have almost identical daily volatility (6.72 vs 5.60 bp) but wildly different cumulative outcomes (-6.16 vs -0.82 pp). The drift-to-noise ratio is roughly seventeen times higher inside the window. So

the question is not which component moves, but **why only scheduled-date price changes persist**. A long-horizon variance-ratio decomposition estimated separately by day type converts this from a description into a measured object, and it requires no decomposition, no endpoint measure, and no model of the term premium — so it survives every identification problem above.

This is the direction I think is most promising and it is the one I would build a chapter around.

Contingent

(d) **Extend the Fed-side measure to the pre-2012 period**, where the effect actually lives, using Tealbook/Greenbook projections (available 1981–2019 with the five-year lag). This is the only way to test any projection-based mechanism on the sample that contains the phenomenon.

7 Where I would value guidance

1. Whether the reconciliation is better as a standalone short submission or folded into a chapter as a methodological section.
2. Access to Blue Chip Financial Forecasts long-range and to the Tealbook archive — both are binding constraints on (d).
3. Whether the permanent/transitory framing in (c) is the right way to pose the question, or whether it is too descriptive to carry a chapter without a mechanism.

Data. Gürkaynak–Sack–Wright zero-coupon yields; Adrian–Crump–Moench daily decomposition, June 1989 – August 2026; FOMC calendar assembled from Federal Reserve historical meeting pages (300 scheduled meetings 1989–2026 plus 73 conference calls and unscheduled actions); SEP dot-plot archive, 57 meetings, January 2012 – March 2026. All estimates recomputed from source; nothing taken from a prior write-up.