

The FOMC Window

What is identified, what is not, and where the project goes next

Progress report · 11 August 2026

1 The motivating fact

The starting point is Hillenbrand (2025, *RFS*). Between June 1989 and June 2021 the ten-year Treasury yield fell 6.99 percentage points. Of that, 6.16 pp — 88% — accrued inside three-day windows around FOMC announcements (days $t - 1$, t , $t + 1$), which are only 11.8% of trading days. On the other 88% of days the yield went essentially nowhere.

I replicate the result from Gürkaynak–Sack–Wright zero-coupon yields and a hand-assembled FOMC calendar (300 scheduled meetings 1989–2026 plus 73 conference calls and unscheduled actions). Figure 1 reproduces it.

1.1 The fact is really an absence

It is worth being precise about what is surprising here, because the per-meeting effect is small: about -2.5 bp. The striking quantity is the *non-window* drift, which is -0.012 bp/day with $t = -0.18$ across 7,035 days. Thirty-two years of macroeconomic news, inflation surprises, recessions and QE announcements net out to zero outside a three-day calendar.

A randomisation test confirms the concentration is not chance: placing 324 randomly timed three-day blocks on non-FOMC days, 5,000 times, puts the observed -6.16 pp at the 0th percentile of the simulated distribution ($p < 0.0002$).




Figure 1: Long rates fall only when the Fed is in the room. Cumulative change in the ten-year yield from 2 June 1989 (8.52%), split into the part accruing inside three-day FOMC windows and the part accruing on all other days. Paper sample, through June 2021.




Figure 2: Same volatility, opposite persistence. Daily standard deviation and absolute mean drift of ten-year yield changes, window days versus all other days, June 1989 – June 2021.

The two day types are not distinguished by volatility. Window days have a daily standard deviation of 6.72 bp against 5.60 bp elsewhere — a 20% difference. They are distinguished by *drift*: -0.634 bp/day against -0.012 . The drift-to-noise ratio is roughly 45 times higher inside the window (Figure 2). Announcement days are not noisier days; they are days on which the noise does not cancel.

1.2 Hillenbrand’s explanation: long-run Fed guidance

Hillenbrand’s own interpretation is that FOMC announcements are the moments at which the Fed reveals information about *where rates are going in the long run* — what he calls **long-run Fed guidance**. The premise is an information asymmetry: by continuously observing how the economy responds to the rates it sets, the Fed accumulates an estimate of the natural rate that markets cannot form as precisely on their own. Investors learn that estimate at announcements and reprice long-duration bonds accordingly. Because the object being revised is an *endpoint* rather than a cyclical position, the revision is permanent, and permanent revisions cumulate rather than cancel — which is exactly the pattern in Figure 1.

He describes the guidance as operating through two channels:

- **Implicit, before 2012.** When the Fed moves the funds rate, the move itself reveals its reading of where the natural rate sits. Consistent with this, the meetings at which the target changed — roughly 45% of meetings in his pre-2012 sample — account for almost the entire decline in long yields over that period.
- **Explicit, from 2012.** The SEP dot plot publishes each participant’s longer-run federal funds projection, communicating the Fed’s endpoint view directly.

Four pieces of evidence are offered that the object being moved is a long-run *real* endpoint rather than the cycle, inflation, or a risk premium.

1. **Far-forward rates.** The pattern holds for the 5-year, 5-year-forward rate, where the

current policy stance has largely washed out. A cyclical explanation should not survive that far out the curve.

2. **Real rather than inflation.** Since 1999 the decline is essentially all in TIPS yields, with breakeven inflation roughly flat — so it is Δr^* , not $\Delta \pi^*$.
3. **The dot plot links the two legs directly.** Regressing window changes in long-term real yields on revisions to the Fed’s longer-run projection, a 100 bp cut in the Fed’s own long-run forecast is associated with more than a 70 bp fall in long-term real yields — a pass-through large enough to be about the endpoint rather than the near path.
4. **Persistence.** The series built only from window changes tracks the low-frequency movement of the ten-year closely, while the non-window series is transitory and offsets over time. Endpoint news should not reverse; cyclical or liquidity-driven news should.

In terms of the identity developed in §2.1, the claim is a strong one: in FOMC windows $\Delta y^{(n)} \approx \Delta r^*$, with the inflation anchor, the policy-stance terms and the term premium all contributing approximately nothing. Each of those three exclusions is a separate assumption, and the rest of this report is largely about which of them survive.

2 The question

If the Fed sets an overnight rate, why does the *ten-year* rate move only on its calendar — and why do those moves not wash out? The question I set out to answer is which component of the yield the Fed is moving:

When the Fed moves long rates at scheduled announcements, is it changing **where markets think rates are going** (expected future short rates), or **what markets charge for the risk of being wrong about that** (the term premium)?

2.1 Unpacking the identity

The reason this leads directly to r^* is an accounting identity, and it is worth writing out in full because the interesting term is the one usually waved away.

Start from the definition of a long yield as an average of expected short rates plus a residual:

$$y_t^{(n)} = \underbrace{\frac{1}{n} \sum_{j=0}^{n-1} E_t[i_{t+j}]}_{\text{expectations component } rn_t^{(n)}} + TP_t^{(n)}. \quad (1)$$

This is a definition, not an assumption: $TP_t^{(n)}$ is whatever is left over.

Now split the nominal policy rate into a neutral level and a deviation from it. Define the *neutral nominal rate* and the *policy gap*

$$i_t^* \equiv r_t^* + \pi_t^*, \quad \text{gap}_t \equiv i_t - i_t^* = \underbrace{(i_t - E_t \pi_{t+1} - r_t^*)}_{\text{real policy stance}} + \underbrace{(E_t \pi_{t+1} - \pi_t^*)}_{\text{inflation deviation}}, \quad (2)$$

where r_t^* is the equilibrium real rate and π_t^* the long-run inflation anchor. The gap is therefore **the stance of policy**: how far the Fed is holding the short rate above or below the level that

would be neutral, plus how far inflation is expected to run from its anchor. Under a Taylor-type rule $i_t = i_t^* + \phi_\pi(\pi_t - \pi_t^*) + \phi_x x_t + v_t$, the gap is exactly the systematic and discretionary response to the cycle,

$$\text{gap}_t = \phi_\pi(\pi_t - \pi_t^*) + \phi_x x_t + v_t. \quad (3)$$

It is a cyclical, mean-reverting object: it is large in 1994, in 2008 and in 2023, and zero on average.

Substituting (2) into (1) and adding the assumption that the endpoints are (locally) martingales, $E_t[r_{t+j}^*] = r_t^*$ and $E_t[\pi_{t+j}^*] = \pi_t^*$, gives the identity used throughout:

$$y_t^{(n)} = r_t^* + \pi_t^* + \frac{1}{n} \sum_{j=0}^{n-1} E_t[\text{gap}_{t+j}] + TP_t^{(n)}. \quad (4)$$

How much does the gap term actually contribute?

The usual move at this point is to say the gap “averages out” at a ten-year horizon. That is too quick, and it is the step I initially got wrong. If the gap mean-reverts as an AR(1) with monthly persistence ρ , the expected-gap term collapses to a single scalar times today’s stance,

$$\frac{1}{n} \sum_{j=0}^{n-1} E_t[\text{gap}_{t+j}] = \omega(\rho, n) \text{gap}_t, \quad \omega(\rho, n) = \frac{1 - \rho^n}{n(1 - \rho)}, \quad (5)$$

and for a forward rate covering periods n_1 through $n_1 + n_2 - 1$ the analogous weight is $\omega^f = \rho^{n_1} (1 - \rho^{n_2}) / [n_2(1 - \rho)]$.

Two objects enter ω : the maturity n , which is known, and the persistence ρ , which is a *belief* and is not directly observed. So the relevant exercise is a sensitivity analysis over ρ , expressed as the half-life of a stance deviation — the time markets expect a 100 bp deviation from neutral to decay to 50 bp:

Half-life of the gap	ω , 2-year	ω , 5-year	ω , 10-year	ω^f , 5y5y forward
1 year	0.56	0.29	0.15	0.01
2 years	0.73	0.48	0.28	0.09
3 years	0.81	0.60	0.39	0.19
4 years	0.85	0.67	0.48	0.28
5 years	0.88	0.73	0.54	0.36

The upper rows are not the empirically relevant ones. Estimating an AR(1) on a crude gap proxy — the one-year yield less (HLW r^* plus 2) — at monthly frequency gives a half-life of 4.4 years on the full sample, 5.0 years post-1985 and 4.1 years post-1995. The bottom rows of the table are where the data sits, so $\omega(10y) \approx 0.5$: a 100 bp revision to the expected stance — with no change whatsoever in r^* , π^* or the term premium — moves the ten-year yield by roughly half as much.

That estimate is an upper bound and should be treated as one. Any slow drift in the true neutral rate that the proxy misses — the disinflation, movements in r^* that HLW does not capture — is dumped into the residual and reads as gap persistence. A better π^* series would pull the half-life down. But it would have to fall to roughly one year before $\omega(10y)$ reached 0.15.

Why this bears on the far-forward evidence

The last column is the one that matters for Hillenbrand’s second piece of evidence (§1.2). Testing the pattern on the 5-year, 5-year-forward rate is exactly the right instinct: the first five years of the expected path drop out, so the stance term is mechanically damped. But the test is only *decisive* if the gap is fast. At a one-year half-life the 5y5y weight is 0.01 and the argument is airtight; at four years it is 0.28, and a 100 bp stance revision still moves the 5y5y forward by nearly 30 bp — comfortably enough to generate the observed window moves without any r^* news at all.

At the persistence actually estimated above, the far-forward result does not separate endpoint news from stance news. Whether it does is an empirical question about ρ that this literature has not addressed, and it is cheap to address: the half-life of the policy gap is estimable, and the sensitivity of the conclusion to it is one line of arithmetic.

The identification problem, stated

Differencing (4) over an announcement window,

$$\Delta y_t^{(n)} = \underbrace{\Delta r_t^* + \Delta \pi_t^*}_{\text{endpoint news}} + \underbrace{\omega \Delta \text{gap}_t + \text{gap}_t \Delta \omega}_{\text{stance and persistence news}} + \Delta TP_t^{(n)}. \quad (6)$$

Four terms, one observable. Note that the second stance term, $\text{gap}_t \Delta \omega$, is a revision to *how long the current stance will last* — which is what forward guidance is, and which is largest precisely when the stance is large.

Three consequences for the project

1. “The Fed moves long rates” does not imply “the Fed moves r^* beliefs.” The stance terms are a live alternative at the persistence the data suggest, and the $\text{gap}_t \Delta \omega$ piece is largest exactly when the stance is large — 2008–09 and 2022–24. It is also plausibly what Bauer, Pflueger & Sunderam’s perceived-reaction-function result is picking up.

2. π^* is not fixed over most of the sample. The claim that the inflation anchor is constant is defensible after roughly 2000 and indefensible before it: the 1989–1999 window declines coincide with the disinflation and with the Fed acquiring credibility, so $\Delta \pi^*$ plausibly carries a large share of them. Since 81% of the in-window decline predates 2012 (§5.3), this is not a small caveat — it means the pre-2012 evidence cannot be read as r^* learning either. TIPS breakevens begin in 1997 (liquid from about 2003), so the decomposition is only testable on the back half.

3. The identity is what makes the term-premium question sharp. If $\Delta \pi^*$ and the stance terms could be shut down, then a permanent window move in the ten-year would have to be Δr^* or ΔTP , and the whole question reduces to splitting those two. That is the design in §5 — and it is why the failure of the market-side endpoint measure is fatal rather than inconvenient.

3 What the literature has established

3.1 The core fact and its extensions

Paper	What it establishes
Hillenbrand (2025, <i>RFS</i>)	The whole secular decline in the ten-year accumulates in three-day FOMC windows, 1989–2021. Attributes it to markets learning a lower r^* from the Fed.
Hofmann, Li & Wu (BIS WP 1252)	The pattern is global — ten advanced economies — and survives a shifting-endpoint model.
Hanson, Lucca & Wright (2021, <i>QJE</i>)	Announcement-window long-rate moves <i>substantially reverse</i> within 6–12 months, which they read as evidence of temporary term-premium effects.
Bauer & Swanson (2023, <i>NBER Macro Annual</i>)	Announcement surprises are predictable from pre-meeting public data; the “Fed information effect” is largely the Fed responding to news.

Rows 1 and 3 are in direct tension. Hillenbrand’s central fact requires that window moves *not* reverse — they have to cumulate into a 32-year trend. Hanson–Lucca–Wright report that they largely do reverse. The two claims have never been confronted directly.

3.2 The contested channel

Paper	What it establishes
Hofmann, Li & Wu (BIS WP 1252)	The window decline is risk-neutral (US pseudo- R^2 96 versus 32). But their shifting endpoint is <i>defined</i> as the first principal component of cumulative FOMC-window yield changes, so the risk-neutral component inherits the series being explained.
Pan & Peng (SSRN 4764451)	On day $t - 1$ the decline is term premium : -0.71 bp of a -0.79 bp move, with expectations -0.08 and insignificant.
Joslin, Singleton & Zhu (2011); Joslin, Pribsch & Singleton (2014)	Under spanning the state is recoverable from yields, so the expectations component is a fixed linear function of the curve.
Bauer & Rudebusch (2020, <i>AER</i>); (2014, <i>IJCB</i>)	A stationary VAR books genuine endpoint shifts as term premium: 75% versus 35–37% for the secular decline.

Two papers, opposite answers, no adjudication. Hillenbrand himself reports in the Internet Appendix that ACM and Kim–Wright term premia fall 100–200 bp in these windows, then sets the finding aside because the models “do not explicitly account for the secular decline” — a reason the measurement is unreliable, not evidence the channel is absent.

3.3 The beliefs-about- r^* strand

Paper	What it establishes
Bauer, Pflueger & Sunderam (2024, <i>QJE</i>)	A pre-meeting <i>subjective belief</i> — the market’s perceived policy reaction-function <i>slope</i> — is related to term premia and to state-dependent announcement-window responses. This is the template for the design.
Ahonon & Roussellet (FRBNY SR 1187)	Under Bayesian learning about unobserved trends, investor uncertainty about r^* is roughly ± 170 bp rather than ± 55 bp, and r^* is the primary long-run driver of bond risk premia.
Couture (2021); Engstrom (FEDS 2026-026)	The dot plot moves asset prices; SEP-minus-market gaps predict forecast errors.
Favero, Melone & Tamoni (2024, <i>JFQA</i>)	A policy-rate-minus-equilibrium-rate gap predicts bond returns. Different wedge, similar name — must be differentiated explicitly.

What has never been tested is the perceived *endpoint*: not how strongly markets think the Fed responds, but where markets think the Fed is heading relative to where the Fed says it is heading.

4 What our extensions add

4.1 Extension 1: the event list decides the answer

Pan & Peng and Hofmann–Li–Wu reach opposite conclusions on the same question. I reproduce Pan & Peng’s estimates on their specification — scheduled meetings only, September 1994 onward, day $t - 1$ — to three coefficients:

Day $t - 1$, ten-year	Mine	Theirs
Yield	-0.785 ($t - 2.37$)	-0.79 ($t - 2.39$)
Term premium	-0.716 ($t - 2.26$)	-0.71 ($t - 2.34$)
Expected short rate	-0.070 ($t - 0.31$)	-0.08 ($t - 0.26$)

and then show the attribution is not robust to the event list. Running the same regression on a 2×2 of {scheduled only, all events} $\times$ {1994+, 1989+} gives:

Sample	n	Total	Expectations	Term premium	TP share
Scheduled, 1994+ (P&P)	250	-0.785	-0.070	-0.716	91%
Scheduled, 1989+	292	-0.755	-0.039	-0.717	95%
All events, 1994+	286	-0.967	-0.415	-0.552	57%
All events, 1989+	364	-0.971	-0.521	-0.450	46%

bp per event; HAC(4) standard errors; sample ends December 2025.

The event list does the work, not the sample start. Holding the list at “scheduled,” moving the start from 1994 to 1989 changes the term-premium share from 91% to 95%. Holding the start at 1994 and adding the unscheduled events changes it from 91% to 57%.




Figure 3: The event list decides the answer. Left: the day- $(t-1)$ decomposition under three event-list definitions — same day, same data, same decomposition. Right: the expectations component per event, scheduled meetings versus unscheduled actions.

The mechanism is identifiable. The 73 conference calls and intermeeting actions that a scheduled-only filter removes carry -1.78 pp of cumulative expectations decline on day $t-1$ alone — about -2.5 bp per event, against -0.07 bp per scheduled meeting, a factor of 35 (Figure 3). Emergency actions are preceded by days of expectations repricing; scheduled announcements are not.

What this adds. The two contradictory papers are reconciled without taking a side on the model. Neither discusses the event list. This is a finished result requiring no model to be correct.

4.2 Extension 2: the sample through July 2026

Extending the panel to 31 July 2026 covers a period no paper in this literature reaches — Hillenbrand ends June 2021, Hofmann–Li–Wu 2023, the Adrian et al. vintage November 2021 — and it is the one period in the sample in which policy tightened sharply.

Between March 2022 and July 2026 the Fed raised the target range 525 bp and revised its own longer-run dot *up* by 0.74 pp. The ten-year rose 2.60 pp. Inside FOMC windows it *fell* 1.20 pp; the entire increase occurred outside (Figure 4). The in-window series continues its decline straight through the tightening, to -7.24 pp.

What this adds. The pattern is not symmetric in the direction of policy, and the symmetry assumption is implicit in every existing treatment. It also means the secular-decline interpretation cannot be the whole story: the in-window decline continued while the Fed was raising its own r^* estimate.

4.3 Extension 3: the decomposition cannot arbitrate

Both contested papers rest on an affine decomposition, so it matters what that decomposition can and cannot deliver at event frequency.

In a spanned affine model the state is recoverable from yields, so the expectations component




Figure 4: The 2022–2026 tightening happened outside the window. Same decomposition as Figure 1, extended to 31 July 2026. Shaded region marks the tightening cycle.

is a fixed linear function of the curve:

$$\Delta rn_t^{(n)} = c'_n \Delta y_t^{\mathcal{N}}, \quad c'_n = b'_n \mathbf{B}^{-1}, \quad (7)$$

where $\mathcal{N}$ is any set of maturities spanning the state. Regressing the daily change in the ACM ten-year risk-neutral rate on six yield changes (1, 2, 3, 5, 7, 10 years) over 9,282 daily observations gives $R^2 = 1.00000$ (Figure 5, left).

The consequence for event studies is immediate: component responses are determined by the maturity profile of the yield response,

$$\beta^{rn} = c' \beta^{\mathcal{N}}, \quad (8)$$

so the decomposition adds *no event-specific information*. Applied to my estimated profile, (8) predicts +2.5 and +29.3 bp against directly estimated +2.6 and +29.2. Running the decomposition is arithmetic on a profile already in hand.

Those weights are also known to be wrong for this application. Since the VAR is stationary, $b_n \rightarrow 0$ and the expectations component is asymptotically constant, so genuine endpoint variation is mechanically booked as term premium. Empirically $c' \mathbf{1} = 0.529$: a permanent, fully credible parallel shift — 100% expectations by construction — is reported as 53% expectations and 47% term premium (Figure 5, right).

Attribution. The algebra in (7) is a corollary of Joslin, Singleton & Zhu (2011) and appears as a criticism in Joslin, Priebisch & Singleton (2014). The stationarity bias is Bauer & Rudebusch (2020, *AER*), who report 75% versus 35–37% for the secular decline, and Bauer & Rudebusch (2014), who report the $\approx 50/50$ event-window figure. My contribution is the explicit filter for ACM and the event-study corollary (8), not the theorem.




Figure 5: The decomposition is a fixed filter of the curve. Left: actual daily changes in the ACM ten-year risk-neutral rate against the fitted values from six contemporaneous yield changes. Right: how a permanent, fully credible 100 bp parallel shift is reported.

5 Where r^* enters

5.1 The design the identity suggests

Equation (4) points at a clean test. If the Fed moves long rates by revealing its own view of the endpoint, then the pre-meeting *gap* between the Fed’s published longer-run projection and the market-implied endpoint should predict the window response — and the sign of the loading on the two components should discriminate the mechanisms. An information channel predicts a signed wedge loading on expectations; a risk channel predicts the magnitude of disagreement loading on the term premium.

This requires two measured objects: a Fed leg and a market leg.

5.2 The treatment exists; the control does not

The Fed leg is real and it moves. The SEP longer-run federal funds projection, less the 2% inflation objective, falls 1.78 pp from 2012 to March 2022 and then rises 0.74 pp (Figure 6, left). The cross-participant mean has 48 distinct values across 57 meetings, against 12 for the median — so the mean, not the median, is the usable series.

The market leg fails. The natural candidate — the ACM 9-year, 1-year-forward risk-neutral rate — is not an endpoint estimate at all. Regressed on the change in the one-year yield it gives $\beta = 0.287$ with $R^2 = 0.83$, and the level correlation is +0.984 (Figure 6, right). It is a damped copy of the current policy cycle.

This is structural rather than a bad horizon choice. Under a stationary VAR the forward loading $b_n \rightarrow 0$, so pushing the horizon out drives the forward toward an unconditional constant, never toward a time-varying endpoint. No horizon works. Survey-based alternatives move too little to help: they correlate -0.12 with the ACM series and barely vary.




Figure 6: An observable treatment and an unmeasurable control. Left: the FOMC’s own r^* , the SEP longer-run projection less 2.0. Right: the ACM-implied “market r^* ,” the one-year yield, and HLW r^* for comparison.

5.3 And the treatment window is the wrong decade

Even setting the measurement problem aside, the design is defeated by timing. Splitting the in-window drift by period:

Period	Meetings	bp/meeting	t	Cumulative
Jun 1989 – Dec 2011 (pre-SEP)	245	−2.46	−3.06	− 5.85 pp
Jan 2012 – Mar 2022 (dot falling 1.78 pp)	85	−0.40	−0.41	−0.23 pp
Mar 2022 – Jul 2026 (dot rising 0.74 pp)	35	−3.49	−1.53	−1.16 pp

The cumulative column sums daily in-window changes; days shared by adjacent events are counted once.

Two things follow. First, **81% of the in-window decline is complete before the dot plot begins** (Figure 7). The period in which the Fed’s endpoint view is observable is the period in which the phenomenon being explained is essentially absent. Every projection-based test of the information channel — including the published one — is run where there is little left to explain.

Second, during the decade when the FOMC cut its own longer-run rate by 1.78 pp — the episode in which an information channel should be loudest — the in-window contribution is −0.40 bp per meeting with $t = -0.41$. Indistinguishable from zero.




Figure 7: The effect and the measurement do not overlap in time. Left: cumulative in-window decline, with the SEP era shaded. Right: in-window drift per meeting by period.

6 What survives

Result	Identifying assumption	Status
Event-list reconciliation: TP share 91% $\rightarrow$ 46%	None beyond the ACM series being the same object in both papers	Finished; most immediately publishable
ACM is a fixed filter, $R^2 = 1$, $c'1 = 0.529$	Spanning; ACM's own estimated loadings	Finished; theorem standard, application new
Timing fact: 81% pre-dates the SEP	Window assignment only	Finished; reframes the projection-based literature
2022–26: tightening occurs entirely outside windows	Window assignment only	Finished; no other paper covers it
Fed–market wedge design	Requires a non-circular market end-point	Set aside — no such measure available

Two further lines were considered and dropped. A credibility asymmetry between dovish and hawkish signals is underpowered (18 upward revisions gives 21–31% power), and Hillenbrand's own Table 5 and Figure 11 already split by direction while the ratchet mechanism is Rungcharoenkitkul, Borio & Disyatat (BIS WP 817). Expectations versus term premium as a standalone question is contested by the two papers above and unidentifiable for the reason in §4.3.

7 The unanswered question I would build on

Return to Figure 2. Window and non-window days have nearly identical daily volatility (6.72 versus 5.60 bp) but completely different cumulative outcomes (−6.16 versus −0.82 pp). The drift-to-noise ratio is roughly 45 times higher inside the window.

So the question is not which component moves. It is:

Why do only scheduled-date price changes persist?

A long-horizon variance-ratio decomposition estimated separately by day type converts this from a description into a measured object: the permanent share of a day’s innovation, estimated for window days and non-window days separately, over horizons of one month to two years. It requires no decomposition, no endpoint measure and no model of the term premium — so it survives every identification problem above. It also confronts Hanson–Lucca–Wright (reversal) with Hillenbrand (permanence) directly, on a statistic that neither reports.

8 Path forward

Immediate (four weeks)

(a) **The issuance test.** Lou, Pintér, Üslü & Walker (BIS WP 1232) document a supply-driven pre-announcement drift in UK gilts and claim Hillenbrand’s effect concentrates in FOMC windows *without* long-dated Treasury issuance. If correct, the mechanism is dealer inventory rather than monetary policy. Auction dates are public; this is a few days’ work and it is decisive either way.

(b) **Draft the reconciliation note** — Extension 1 plus the timing fact in §5.3. Self-contained, model-free, and it resolves a live disagreement.

Medium term (one term)

(c) **Permanent versus transitory by day type** — the variance-ratio design in §7. This is the direction I think is most promising and the one I would build a chapter around.

Contingent

(d) **Extend the Fed-side measure back before 2012**, where the effect actually lives, using Tealbook/Greenbook projections (available 1981–2019 with the five-year lag). This is the only way to test any projection-based mechanism on the sample that contains the phenomenon.

9 Where I would value guidance

1. Whether the reconciliation is better as a standalone short submission or folded into a chapter as a methodological section.
2. Access to Blue Chip Financial Forecasts long-range and to the Tealbook archive — both are binding constraints on (d).
3. Whether the permanent/transitory framing in (c) is the right way to pose the question, or whether it is too descriptive to carry a chapter without a mechanism.

Data. Gürkaynak–Sack–Wright zero-coupon yields; Adrian–Crump–Moench daily decomposition, June 1989 – August 2026; FOMC calendar assembled from Federal Reserve historical meeting pages (300 scheduled meetings 1989–2026 plus 73 conference calls and unscheduled actions); SEP dot-plot archive, 57 meetings, January 2012 – March 2026. All estimates recomputed from source; nothing taken from a prior write-up.